Distributed economic dispatch of embedded generation in smart grids

J. Dibangoye [a], A. Doniec [b,*], H. Fakham [c], F. Colas [c], X. Guillaud [d]

[a] INSA-Lyon, CITI-INRIA, F-69621, France
[b] URIA – Ecole des Mines de Douai, France
[c] L2EP – Arts et Métiers ParisTech, France
[d] L2EP – Ecole Centrale de Lille, France

A B S T R A C T

In a smart grid context, the increasing penetration of embedded generation units leads to a greater complexity in the management of production units. In this paper, we focus on the impact of the introduction of decentralized generation for the unit commitment (UC) problem. Unit commitment problems consist in finding the optimal schedules and amounts of power to be generated by a set of generating units in response to an electricity demand forecast. While this problem has received a significant amount of attention, classical approaches assume that these problems are centralized and deterministic. However, these two assumptions are not realistic in a smart grid context. Indeed, finding the optimal schedules and amounts of power to be generated by multiple distributed generator units is not trivial since it requires to deal with distributed computation, privacy, stochastic planning, etc. In this paper, we focus on smart grid scenarios where the main source of complexity comes from the proliferation of distributed generating units. In solving this issue, we consider distributed stochastic unit commitment problems. We introduce a novel distributed gradient descent algorithm which allows us to circumvent classical assumptions. This algorithm is evaluated through a set of experiments on real-time power grid simulator.

© 2015 Elsevier Ltd. All rights reserved.

1. Introduction

The economic dispatch and unit commitment (UC) problem consists in finding the optimal schedules and amounts of power to be generated by a set of power generators (units) in response to an electricity demand over a planning horizon (Aoki et al., 1989; Borghetti et al., 2001; Guan et al., 2003). Earlier approaches for solving UCs including branch-and-bound methods, dynamic programming and Lagrangian relaxation techniques assume that units are fully reliable and share all together their states, technical specifications and schedules (Cohen and Yoshimura, 1983; Snyder et al., 1987; Fisher, 2004). However, the increasing penetration of embedded units in distributed networks together with the liberalization of electricity markets make these assumptions less and less realistic on both demand and supply sides (Kok et al., 2010; Nikovski and Zhang, 2010; Ramchurn et al., 2012).

On the demand side, more and more customers supplement the amount of power their own units generate by that of the electrical utilities, which makes demand forecast inaccurate. On the supply side, the amount of power generated by an electrical utility influences the amount of power other electrical utilities need to generate in order to meet the demand. Furthermore, the liberalization of the electricity markets precludes electrical utilities to share their private information with one another including: schedules, generation capabilities, technical specifications, generator failure histories, blackouts, etc. As a consequence, centralized and deterministic models are no more relevant to unit commitment problems. Even more importantly, these limitations highlight the impetus for models that can produce operational schedules that are robust in face of both: supply and demand uncertainties and privacy-preserving constraints.

Traditional responses to supply and demand uncertainties have been to schedule enough reserve so as to face forecast inaccuracies or generator failures. Typically, a safety margin of 3% in the production is commonly used in power generation as a *reliability rule-of-thumb* (Sheble and Fahd, 1994). This heuristic strategy often results in the generation of amounts of power that significantly exceed the expected demand, and thus the operational costs of electrical utilities are overestimated. Clearly, as the penetration of embedded units increases, such heuristics are likely to overestimate the operational costs and amounts of power units generated.

A more promising approach assumes that the uncertainty constraints are parts of unit commitment problems, which makes

the latter stochastic. The goal, then, consists in finding schedule strategies that minimize the expected operational costs while preserving the ability to meet the expected demand, and ensuring the robustness in face of supply and demand variability. Schedule strategies implicitly provide safety margins by taking into account all possible contingencies. Notice that the idea of using stochastic unit commitment (SUC) problems in order to deal with supply and demand uncertainties is not new. It can be traced back to Takriti et al. (1996), who developed a stochastic programming model and solution method based on Lagrangian relaxation techniques.

Since then, numerous authors have refined both the model and the solution, exploiting Lagrangian heuristics (Nowak and Rmisch, 2000), security-based probabilistic models (Bouffard et al., 2005), market-based mechanisms (Vytelingum et al., 2010) and Markov decision processes (Nikovski and Zhang, 2010) to cite a few. Unfortunately, the number of all possible contingencies in SUCs may grow exponentially with the planning horizon, making exact approaches intractable (Nikovski and Zhang, 2010). Instead of considering the entire contingencies, Takriti et al. (1996) suggest to plan only over a few scenarios, which significantly improves the scalability of the solution method. But there is no free lunch, such an approach often fails to address all future possible realizations that are not part of the selected scenarios.

Though approaches to solving SUCs can handle uncertainty, they all assume electrical utilities share with one another all their private information. That is, there exists a centralized coordinator agent that computes a centralized schedule strategy on behalf of the entire set of electrical utilities. However, the liberalization of electricity markets tends to enforce a system of competition where electrical utilities compete to offer their electricity output to retailers, making centralized approaches no longer reliable. In such a setting, schedule strategies, the centralized coordinator agent computes, are obsolete as they centralize private information of all electrical utilities.

To tackle the privacy-preserving bottleneck, the past few years have seen many distributed approaches to preserve private information of electrical utilities involved in a distributed system. In distributed approaches, each electrical utility is an autonomous processing node, we will call an *agent*, which works together with the other nodes in order to solve a unit commitment problem. The agents collaborate to coordinate their resources and activities while preserving their private information. Notable examples include the work by Kim and Baldick (1997), who developed a distributed algorithm that extends deterministic and centralized Lagrangian relaxation methods; or that of Miller et al. (2012), who introduced a message passing algorithm to UCs in the form of distributed constraint optimization problems (Kumar et al., 2009; Modi et al., 2005). Unfortunately, none of these distributed approaches can handle the uncertainty in supply and demand. So, it would seem like we are constrained to either face the variability of supply and demand, or preserve private information of electrical utilities. To the best of our knowledge, none of the existing approaches can overcome both: supply and demand uncertainties; and privacy-preserving constraints.

In this paper, we introduce an algorithmic framework that extends both stochastic and distributed approaches to UCs in order to ensure that electrical utilities do not explicitly communicate their private information to their competitors during the planning phase. In particular, we recast distributed stochastic unit commitment (DSUC) problems into linearly constrained quadratic programs (LCQP). In this form, the primary contribution of this work is to extend existing distributed algorithms for solving unconstrained quadratic programs to LCQP and thus DSUC. This is achieved by means of communication protocols that allow electrical utilities to choose which part of their private information to share with one another in order to collectively find an optimal or

near-optimal schedule strategy. The resulting algorithm, namely *protocol based distributed projected gradient-descent optimization* (P-DPGO), is guaranteed to terminate after a finite number of iterations with a near-optimal solution.

We demonstrate the performance of the P-DPGO algorithm on an IEEE 14 nodes network, which consists of several virtual power plants with controllable generator units. The principal source of uncertainty in such a setting is the unpredictable break-downs generator units can experience, which make the supply uncertain. Experiments over different stochastic scenarios show that P-DPGO produces efficient schedule strategies in term of costs. This is expected given that our approach exploits three advantages: first, it preserves private information of electrical utilities; next, it takes into account long term decision effects; and finally, it can handle uncertainty in supply and demand.

The remainder of this paper is organized as follows. First, we provide some motivating scenarios that illustrate the key features in DSUCs (Section 2). Next, Section 3 describes different models of UCs, and Section 4 discusses distributed algorithms for solving unconstrained quadratic programs we build upon. Then, in Section 5, we describe P-DPGO, which combines existing distributed algorithms to communication protocols in order to ensure that local information electrical utilities do not want to share remain private. Finally, we present an empirical evaluation of this algorithm on a real experimental platform.

2. Motivating scenarios

In the following, we distinguish between three scenarios that illustrate the characteristics of unit commitment problems we target. The primary scenario involves no competition at all and no uncertainty, the second augments the former by taking into account uncertainty and the last scenario complements the second one assuming a competitive setting.

Scenario 1. In this first scenario, we consider a smart grid that consists of two controllable units, e.g., diesel power generators, each of which is owned by a single electrical utility. This electrical utility needs to find the least-cost dispatch of available generation resources to meet the electrical load over 24 h. In this world, each unit can generate power profiles that range from 10% to 95% of its nominal generation capacity. In addition, each unit incurs operations and maintenance costs that increase quadratically with the amount of power it generates. Furthermore, each unit is subject to a number of complex and private technical constraints, e.g., the maximum rate of ramping up or down and the minimum period the unit is up and or down.

Such a non-competitive and deterministic scenario is amenable to centralized branch-and-bound methods, dynamic programming and Lagrangian relaxation techniques (Cohen and Yoshimura, 1983; Snyder et al., 1987; Fisher, 2004). However, there are various sources of uncertainty in real-world unit commitment problems. Examples, such as unpredictable failures in the transmission network or the introduction of uncontrollable units, make consumers at different nodes in the distribution network to experience cuts in the power supplied. As a consequence, the other distribution networks compete to offer their electricity output in order to meet the demand of these consumers. A stochastic scenario, which complements that in Scenario 1, follows.

Scenario 2. In this second scenario, each unit can experience unpredictable breakdowns, which result in cuts in the power supplied. More precisely, at the end of each time step, each unit can sense some information, which in this case corresponds to whether or not a breakdown occurs.

As the breakdowns are not predictable with certainty, the first scenario is no longer reliable. The second scenario is amenable to stochastic models that are robust in the face of the supply variability (Sheble and Fahd, 1994; Takriti et al., 1996). Nevertheless, the second scenario still lacks a key feature in nowadays unit commitment problems. With the liberalization of the electricity market, the power generation sector is organized on a competitive basis with independent electrical utilities selling their production to distribution companies. In order to address the issues of embedded units, the third scenario slightly modifies the second one to include the need to preserve private information of each electrical utility.

Scenario 3. This third scenario modifies the second scenario and assumes a different electrical utility owns each generator unit. Together these electrical utilities need to collaborate (e.g., communicate their schedules) in order to find the least-cost dispatch of available generation resources to meet the electrical load over 24 h. However, each electrical utility is self-interested, i.e., schedules they share with the other electrical utilities focus on individual desires rather than collective desires. As a result, the new electrical market enforces competition among electrical utilities, which precludes them to communicate their private information.

Scenario 3 highlights the necessity to preserve private information that makes centralized solutions obsolete. In the third scenario, if electrical utilities rely on a centralized solution, then they would need to share with one another all their private information, which is not desirable. This highlights the impetus of distributed stochastic models, thereby each unit maintains its private information about the state of the world and chooses which part of this information to share with the others. These characteristics appear in many real-world applications including multiple space exploration rovers used by NASA to explore the surface of Mars and sensor net domains where a team of stationary and moveable UAVs, satellites or other sensors must coordinate to track targets (Zilberstein et al., 2002; Lesser et al., 2003). In the following, we answer the question how these scenarios can be formalized and solved.

3. Formal models

In this section, we introduce models we use to formalize and solve unit commitment scenarios presented in Section 2. We start with a brief review of UC (unit commitment) problems (Section 3.1), then we discuss in Section 3.2 SUCs (stochastic UCs) and finally introduce DSUCs (Distributed SUCs) in Section 3.3.

3.1. Unit commitment

Given load profile (e.g., electricity demand for each hour of a day), a set of units, and operational constraints, the unit commitment problem is concerned with two interdependent goals. First, it decides which units should be started or stopped. Second, it determines the power output of each unit, which will minimize the overall cost needed to meet the load. A formal definition, adapted from Nikovski and Zhang (2010), follows (Table 1).

Definition 1. A unit commitment problem is given by (N, T, X, P, D, C) where

- N is the number of controllable units involved in the system.
- T is the number of periods in the planning horizon, typically a period corresponds to 1 h and the planning horizon corresponds to one day.

Table 1
Notation summary.

Notation	Meaning
N	Number of controllable units
K	Number of uncontrollable units
T	Number of periods in the planning horizon
X, X_i	Joint and unit status domains
P, P_i	Joint and controllable unit output domains
D	Load-demand domain
C, C_j	Joint and unit constraint domains
Y, Y_k	Joint and uncontrollable unit output domains
Z	Joint observation domain
$\underline{A}_i, \overline{A}_i, A_i$	Set of agents with signal $-1, +1$ and 0, respectively
x, x_i	Joint and controllable unit commitment status profiles
p, p_i	Joint and controllable unit output profiles
d	Demand profile
y, y_k	Joint and uncontrollable unit output profiles
z	Joint observation
π, π_i	Joint and unit policies of all units
$\overline{c}, \overline{c}_i, c_i$	Joint and controllable unit hard and soft constraints
$v(\pi)$	Expected value of policy π

- $X \equiv \times_{i \in \{1,2,\ldots,N\}} X_i$ is a joint status domains, where X_i is a status domain for unit i. Let $x \equiv (x(t))_{t \in \{1,2,\ldots,T\}}$ and $x(t) \equiv (x_i(t))_{i \in \{1,2,\ldots,N\}}$ be the commitment status of all units at the t-th period, and $x_i(t) \in X_i$ be the commitment status of unit i at period of time $t \in \{1, 2, \ldots, T\}$.

- $P \equiv \times_{i \in \{1,2,\ldots,N\}} P_i$ is a joint output domain, where P_i is an output domain for unit i. Let $p \equiv (p(t))_{t \in \{1,2,\ldots,T\}}$ and $p(t) \equiv (p_i(t))_{i \in \{1,2,\ldots,N\}}$ to be the output profile in P of all units at the t-th period, and $p_i(t) \in P_i$ be the commitment status of unit i at period of time $t \in \{1, 2, \ldots, T\}$.

- D is the load-demand domain at any period in the planning horizon. We denote $d \equiv (d(t))_{t \in \{1,2,\ldots,T\}}$ to be the electricity demand profile over T periods, and $d(t) \in D$ be the load demand at the t-th period.

- $C \equiv \cup_{j \in \{0,1,\ldots,N\}} C_j$ is a set of constraints that affect either each unit individually (i.e., unit constraints $(C_i)_{i \in \{1,2,\ldots,N\}}$) or the more than one unit (i.e., system constraints denoted by C_0).

This definition offers a high-level description about parameters involved in a unit commitment problem. It is worth providing examples to illustrate some of these parameters. For instance, we need to distinguish between unit i's constraints C_i and system constraints C_0. Unit constraints are constraints that affect only a single unit. We distinguish between: the operational cost function $c_i(x_i, p_i)$, where x_i and p_i are commitment status and load output of unit i; and other units constraints such as the output profile limits $\overline{c}_i(x_i, p_i)$ that are inequality constraints. In contrast, the system constraints affect more than one unit, examples include the load and generation balance, denoted $\overline{c}(x, p)$, and given by $\sum_i x_i(t) \cdot p_i(t) = d(t)$ for all period $t \in \{1, 2, \ldots, T\}$, where x and p are commitment status and output profiles of all units, respectively. It will prove useful to notice that the unit commitment problem involves both equality (e.g., the load and generation balance) and inequality (e.g., each unit's output limits) constraints. For a thorough discussion on unit and system constraints involved in unit commitment problems, the reader can refer to Nikovski and Zhang (2010).

Given a unit commitment problem (Definition 1), the goal is threefold:

1. to set the joint status $x(t) \equiv (x_i(t))_{i \in \{1,2,\ldots,N\}}$ of all units at every period of time $t \in \{1, 2, \ldots, T\}$, where $x_i(t) \in \{0, 1\}$ is the commitment status of unit i at the t-th period of time;

2. to find the joint output $p(t) \equiv (p_i(t))_{i \in \{1,2,...,N\}}$ of all units at every period of time $t \in \{1, 2, ..., T\}$, where $p_i(t) \in P_i$ is the output of unit i at the t-th period of time;
3. to ensure joint output $p \equiv (p(t))_{t \in \{1,2,...,T\}}$ and joint status $x \equiv (x(t))_{t \in \{1,2,...,T\}}$ minimize the overall cost to meet the load $d(t) \in D$ and satisfy other constraints in C at each period of time $t \in \{1, 2, ..., T\}$.

Overall, an optimal operational schedule (x^*, p^*) is the solution of the following quadratically constrained optimization program:

$$(\mathcal{P}) : \begin{cases} \arg\min_{x,p} & \sum_{i=1}^{N} c_i(x_i, p_i) \\ \text{subject to :} & \overline{c}(x, p) \text{ and } \overline{c}_i(x_i, p_i) \quad \forall i \in \{1, 2, ..., N\}. \end{cases}$$

Next, we introduce an illustrative example of unit commitment problem $(\mathcal{P})$.

Example 1. Consider two plants over one-step period, each plant of which having a single controllable generating unit, i.e., $T=1$ and $N=2$. Given load demand $d=200$ MW, our goal is to find generation profiles p_1 and p_2, which minimizes the cumulated quadratic costs:

$\arg\min_{p_1, p_2} c_1(p_1) + c_2(p_2),$
$\quad$ subject to $\quad 200 - p_1 - p_2 = 0,$
$\quad\quad 20 \leq p_1 \leq 125,$
$\quad\quad 39 \leq p_2 \leq 150,$
$\quad\quad c_1(p_1) = 0.5p_1^2 + 215p_1 + 5000,$
$\quad\quad c_2(p_2) = 0.7p_2^2 + 160p_2 + 9000.$

In solving this unit commitment instance, one can use Lagrangian relaxation, which results in $p_1 = 93.75$ and $p_2 = 106.25$ MW. Notice that we do not need the commitment status x_1 and x_2 in such a simple example.

In principle, there exists a number of techniques that can solve problem illustrated in Example 1, including centralized branch-and-bound methods, centralized constraint optimization[1], dynamic programming, and Lagrangian relaxation. For further details on these approaches, the reader can refer respectively to Cohen and Yoshimura (1983), Snyder et al. (1987), Dechter (2003), and Fisher (2004). The following section considers a more challenging scenario, which assumes that UCs are stochastic.

3.2. Stochastic unit commitment

The UC (Definition 1) assumes that all units, load and generation profiles are fully reliable. To cope with various sources of uncertainties including supply and demand variability, we now discuss SUCs, where demand and supply depend on stochastic processes.

Definition 2. A stochastic UC is given by $(K, N, T, X, P, D, C, Y, Z)$ where

- (N, T, X, P, D, C) are similar to UC (Definition 1), except that now $d \in D$ is a random variable drawn from a known stochastic process.
- K is the number of uncontrollable units involved in the system.
- $Y \equiv \times_{k \in \{1,2,...,K\}} Y_k$ is a joint output domain of all uncontrollable units, where Y_k is the output domain of uncontrollable unit k. We denote $y \equiv (y(t))_{t \in \{1,2,...,T\}}$ the joint output profile and $y(t) \equiv (y_k(t))_{k \in \{1,2,...,K\}}$ the joint output at the t-th period, where $y_k(t)$ is the output of uncontrollable unit $k \in \{1, 2, ..., K\}$ at period of time $t \in \{1, 2, ..., T\}$.

- Z is an observation domain (e.g., information about a unit breakdown, a failure at a distribution network node, etc.) about the system the units receive during the planning horizon, where $z \equiv (z(t))_{t \in \{1,2,...,T\}}$ denotes an observation history drawn from a known stochastic process the units receive all together during the planning horizon.

In this paper, we assume that $d(t) \in D$, $y(t) \in Y$ and $z(t) \in Z$ are independent (random) variables, for all periods $t \in \{1, 2, ..., T\}$. Given a SUC, the goal is similar to that of UC (Definition 1), except that operational schedule of all controllable units becomes a policy, denoted π. That is a contingency schedule that prescribes at each period the power output of the entire set of controllable units $\pi(y, z, d) = (x, p)$ conditional on uncontrollable unit output y, observation z and load d profiles. Define the total expected cost of selecting policy π to be

$$v(\pi) = \mathbb{E}_{y,z,d}^{\pi} \left[\sum_{i=1}^{N} c_i(x_i, p_i) | (x, p) = \pi(y, z, d) \right], \tag{1}$$

where $\mathbb{E}$ denotes the expectation operator. Thus, the optimal centralized policy π^* is the solution of the following stochastic quadratically constrained optimization program:

$$(\mathcal{P}') : \begin{cases} \arg\min_{\pi} & v(\pi), \\ \text{subject to :} & \overline{c}(\pi(y, z, d)) \quad \forall y \in Y^T, \ \forall z \in Z^T, \ \forall d \in D^T, \\ & \overline{c}_i(\pi(y, z, d)) \quad \forall i, \forall y \in Y^T, \ \forall z \in Z^T, \ \forall d \in D^T, \end{cases}$$

where $\overline{c}$ and $(\overline{c}_i)_{i \in 1,2,...,N}$ denote system and unit constraints, respectively. Next, we present an illustrative example of a SUC, which extends Example 1 to incorporate the demand variability.

Example 2. Back to Example 1, we now consider that our problem has a random load demand d, where $d=200$ or $d=180$ with equal probability. The goal is to find policy $\pi^* \equiv \{p_{1,d}, p_{2,d} | \ \forall d \in \{200, 180\}\}$, which minimizes the expected quadratically constrained cost:

$\arg\min_{\pi} \frac{1}{2}(c_1(p_{1,200}) + c_2(p_{2,200}) + c_1(p_{1,180}) + c_2(p_{2,180})),$
$\quad$ subject to $\quad 200 - p_{1,200} - p_{2,200} = 0,$
$\quad\quad 180 - p_{1,180} - p_{2,180} = 0,$
$\quad\quad 20 \leq p_{1,200} \leq 125,$
$\quad\quad 20 \leq p_{1,180} \leq 125,$
$\quad\quad 39 \leq p_{2,200} \leq 150,$
$\quad\quad 39 \leq p_{2,180} \leq 150,$
$\quad\quad c_1(p_1) = 0.5p_1^2 + 215p_1 + 5000,$
$\quad\quad c_2(p_2) = 0.7p_2^2 + 160p_2 + 9000.$

Once again, we use Lagrangian relaxation to solve this problem, which results in $p_{1,200} = 93.75$, $p_{2,200} = 106.25$, $p_{1,180} = 82.08$ and $p_{2,180} = 97.92$ MW. Notice that, for the sake of simplicity, we assume that this example involves no observation variables $z \in Z$ and no uncontrollable units.

This example shows what makes stochastic unit commitment problems fundamentally different from deterministic ones. The main difference lies in the number of contingencies, e.g., in Example 2 we have two alternative demand profiles whereas Example 1 considers only one. As the number of contingencies increases the quadratic optimization problem becomes larger, which makes it difficult to find an optimal solution. Nonetheless, a number of solution methods have been developed to solve this problem including stochastic programming based on Lagrangian relaxation techniques (Takriti et al., 1996; Nowak and Rmisch, 2000) and Markov decision process methods (Nikovski and Zhang, 2010) and possibly stochastic constraint optimization

---

[1] Notice that constraint optimization methods applied assuming output domains are discrete and quadratic constrains are made explicit.

methods[2] (Tarim et al., 2006). More specifically, the quadratic optimization problem (Example 2) with equality and inequality constraints can be solved by the method of Lagrangian multipliers.

3.3. Distributed stochastic unit commitment

The stochastic unit commitment problem (Definition 2) assumes units share all together their private information with one another. However, in distributed settings of a stochastic unit commitment problem, each unit is unaware of constraints and cost functions from the other units.

Definition 3. The distributed stochastic unit commitment (DSUC) problem is described similarly to SUC (Definition 2) using tuple $(K, N, T, X, P, D, C, Y, Z)$. The difference lies in the fact that unit constraints $C_i \subset C$ is no longer available to unit $j \neq i$, for all units $i, j \in \{1, 2, \ldots, N\}$.

As no controllable unit has access to the entire stochastic unit commitment problem (Definition 2), it is unlikely that any of them can individually compute the optimal centralized policy π^* on behalf of the entire set of all controllable units. Instead, all controllable units need to collaborate in order to jointly determine π^*. This collaboration takes the form of communication protocols between units, which naturally leads us to distributed techniques. Distribution problem-solving is a well known paradigm with various principled methods. Unfortunately, not all distributed solution methods can solve any distributed problems. For example, distributed methods to solving constraint optimization problems (COPs) (Zhang et al., 2005) are not geared to optimally solving DSUCs. Indeed, COPs and DSUCs formalize two fundamentally different problems: the former is a discrete optimization problem; where the latter is a continuous optimization problem. While it is possible to constrain continuous variables to take their values only in a discrete domain, the final solution has no guarantee to be optimal with respect to the original problem. Overall, to the best of our knowledge, no existing distributed approaches can solve DSUCs. However, there exists distributed techniques that can solve unconstrained quadratic optimization problems. In the following, we progressively build upon these techniques, by integrating all type of constraints involved in distributed unit commitment problems, i.e., equality and inequality constraints.

4. Solving quadratic optimization problems

In this section, we present our distributed solution method to solving a DSUC (Definition 3). To this end, we present solution methods for certain relaxations of DSUCs. In Section 4.1, the first relaxation assumes that all candidate policies satisfy all constraints in C, which makes the resulting problem unconstrained. Next, we present in Section 4.2 another relaxation that assumes all candidate policies satisfy all but equality constraints. Finally, we consider all inequality and equality constraints to address the original problem.

4.1. Solving DSUCs with no constraints

In this section, we review algorithms that have been introduced to solving distributed optimization problems with no constraints. Of this family, the *distributed gradient-descent based optimization* (DGO) algorithm is guaranteed to find an optimal solution to quadratic optimization problems (Tsitsiklis et al., 1986; Mathews et al., 2009). The DGO algorithm is a cooperative distributed problem solving method. That is, a network of autonomous processing agents[3] (i.e., generating units) working together to solve a problem, typically a multi-agent system. In this approach, each agent controls a subset of the decision variables.

At the beginning of each round of the algorithm, agents communicate copies of values assigned to their own decision variables with one another. Then, each agent modifies its own decision variables while keeping fixed values assigned to decision variables of the other agents. In modifying their own decision variables, each agent solves a smaller optimization problem, where solely its own decision variables are unknown. The algorithm continues until two consecutive values assigned to all decision variables are identical, in that case the last assignment is guaranteed to be optimal due to the convexity of the objective function (Tsitsiklis et al., 1986). The optimality guarantees hold in our setting mainly because the objective function is a quadratic function, thus a convex function. In the following, we provide a formal description of DGO and discuss its limitations with respect to equality and inequality constraints.

Before proceeding any further, it is worth noticing that our unconstrained optimization problem $(\mathcal{P}'')$ consists in finding an optimal policy π^* given by

$$(\mathcal{P}'') : \arg\min_\pi v(\pi),$$

where $v(\pi)$ is a quadratic objective function (Eq. (1)), and π describes the values assigned to decision variables of all agents. The distribution of problem $(\mathcal{P}'')$ over all agents results in sub-problem $(\mathcal{P}_i'')$ for each agent $i \in \{1, 2, \ldots, N\}$ given by $(\mathcal{P}_i'') : \arg\min_{\pi_i} v(\pi_i)$, where π_i denotes the values assigned to decision variables of agent i and copies of values assigned to decision variables of the other agents. Let $\pi_i(\tau)$ be the copy of policy π_i that agent i transfers to the other agents at the τ-th round of the DGO algorithm.

All agents start with the same initial solution, possibly through an initial coordination round of the algorithm, i.e., $\pi_1(0) = \pi_2(0) = \ldots = \pi_N(0)$. Then at the τ-th round of the DGO algorithm, each agent i updates $\pi_i(\tau)$ using a gradient-descent approach to solving $(\mathcal{P}_i'')$ – the reader can refer to Boyd and Vandenberghe (2004) for a reference on gradient-descent algorithms. Next, agent i shares copy of $\pi_i(\tau)$ with the other agents through explicit communications, and receives copies $\pi_j(\tau)$ from the other agents $j \neq i$, which in turn will help agent i to update policy $\pi_i(\tau + 1)$, and so on and so forth. Agents stop whenever two consecutive policies are identical.

When applied to unconstrained version of Example 2 such an algorithm is guaranteed to find an optimal solution. Since each unit holds its own cost function and the expected total cost is the addition of private costs, each unit can optimize its private cost on its own. Hence, the first application of a gradient-descent method would result in the optimal policy for each unit.

In the following, we consider a more challenging problem closer to that of Example 2. In this problem, we complement the previous unconstrained problem with equality constraints. This permits to keep local copies of each agent consistent with the equality constraints throughout the distributed planning process.

4.2. Solving DSUCs with equality constraints

This section extends of the DGO algorithm to solving quadratically constrained optimization problems, where constraints are all equality constraints. The resulting algorithm, we call *distributed projected gradient-descent optimization* (DPGO) algorithm, extends

---

[2] The stochastic contraint optimization extends constraint satisfaction or optimization frameworks to deal with uncertainty.

[3] In the following, we use indifferently agent or unit since one agent is in charge of controlling one unit.

the DGO algorithm to deal with equality constraints in distributed stochastic unit commitment problems given by $\forall i \in \{1, 2, \ldots, N\}$,

$$(\mathcal{P}_i''') : \begin{cases} \arg\min_{\pi_i} & v(\pi_i), \\ \text{subject to} & \overline{c}(\pi_i(y, z, d)) \quad \forall y, \forall z, \forall d, \end{cases}$$

where $\overline{c}$ represents the load and generation balance, i.e., an equality system constraint involved in unit commitment problems.

Both DGO and DPGO algorithms share the same algorithmic framework, thereby agents communicate copies of their local solutions until no more changes occur. However, to deal with the equality constraints, agents that need to guarantee their local solutions (i.e., copies owned by each agent) are always consistent with equality constraints. To achieve that, the DPGO algorithm replaces the standard gradient-descent method used in DGO by the *projected gradient-descent method* (Boyd and Vandenberghe, 2004). The projected gradient-descent method is a gradient-descent method where each intermediate solution is projected onto the space of admissible solutions. Hence, the projected gradient-descent method ensures that local solutions are always consistent with the equality constraints.

It is worth noticing that we have not yet considered inequality constraints in the solution of $(\mathcal{P}_i''')_{i \in \{1,2,\ldots,N\}}$. Fortunately, if the solution obtained without considering the inequality constraints satisfies the inequality constraints as well, then the obtained solution will be optimal. If for one or more generator units, the inequality constraints are not satisfied, the optimal policy is obtained by keeping these generator units in their nearest limits and making the other generator units to supply the remaining load demand. However, this would require agents to share their constraints with one another, which is not possible in the new electricity market. To allow agents to take into account these inequality constraints, we introduce communication protocols that permit them to coordinate while preserving private their own information.

5. Solving distributed stochastic unit commitment

This section introduces communication protocols commonly used by agents to ensure that they can find an optimal solution satisfying both system and unit constraints, without actually revealing unit constraints of one generating unit to the others. A straightforward approach would be that each agent can try to satisfy its unit constraints independently from its teammates. Unfortunately, in such a case local solutions from one agent to another may be contradicting or conflicting. Thus, agents may never converge into the same solution. To overcome contradicting or conflicting local solutions, we now introduce communication protocols between agents. These communication protocols are used on the top of the DPGO algorithm, the resulting algorithm is referred to as the *protocol-based distributed projected gradient-descent optimization* (P-DPGO) algorithm.

5.1. Termination conditions

Here we characterize termination conditions of the P-DPGO algorithm when dealing with both equality and inequality constraints. Informally, a termination condition is the set of conditions that are sufficient to ensure that an agent can stop without loosing the ability that the whole team of agents ends with the optimal policy satisfying all constraints. Hence, in identifying termination conditions, we need to distinguish between all scenarios in which an agent can stop and make the whole team of agents lose the ability to find the optimal policy. Before proceeding any further, we start with useful definitions. Specifically, we define what agents communicate and receive.

Definition 4. At each round τ of the algorithm, each agent i *communicates*:

1. a copy of its local solution $\pi_i(\tau)$;
2. a signal $\sigma_i(\tau) \in \{-1, 0, +1\}$, where 0 means all unit constraints are *passive*, $+1$ means one unit constraint has attained its upper-bound limit, and -1 means one unit constraint has attained its lower-bound limit;
3. and the expected cost (bid) $v(\pi_i(\tau))$ it maps to its local solution $\pi_i(\tau)$.

Definition 5. At each round $\tau + 1$ of the algorithm, each agent i *receives*:

1. a copy of local solutions $(\pi_j(\tau))_{j \neq i}$ of the other agents;
2. signals from all agents $j \neq i$ that are spread into two sets $\overline{A}_i(\tau)$ and $\underline{A}_i(\tau)$, where $\overline{A}_i(\tau)$ groups all agents $j \neq i$ that have sent signal $\sigma_j = +1$ and $\underline{A}_i(\tau)$ those that have sent signal $\sigma_j = -1$, the remaining agents are in A_i.
3. and all expected costs $(v(\pi_j(\tau)))_{j \neq i}$ from all agents $j \neq i$.

Notice that agents are allowed to lie since they are in a competitive setting. In such a case, the optimality is with respect to what they communicate to one another.

In the remainder of this subsection, we enumerate four termination conditions (see Theorem 1) that are sufficient to ensure agent termination while preserving the ability to find an optimal solution with respect to available information. To better understand the rational behind these conditions, we show they hold no matter the scenarios the algorithm may face. We distinguish between eight scenarios.[4] A scenario consists in whether or not one of the three sets $A_i(\tau)$, $\underline{A}_i(\tau)$ and $\overline{A}_i(\tau)$ is empty. Note that for each termination condition we consider only a subset of scenarios that fits within the termination condition. A formal characterization of the termination condition and *terminal agents* follows. Agents that are not terminal are referred to *non-terminal agents*.

Theorem 1. *Agent i is terminal iff one of the following conditions hold:*

$\mathcal{C}_1$. *agent i declares its upper-bound limit is active, i.e., $i \in \overline{A}_i(\tau)$;*
$\mathcal{C}_2$. *agent i declares its lower-bound limit is active, part of the other agents declare to satisfy their unit constraints and none of them declare their upper-bounds are active, i.e., $i \in \underline{A}_i(\tau)$, $A_i(\tau) \neq \varnothing$ and $\overline{A}_i(\tau) = \varnothing$;*
$\mathcal{C}_3$. *all agents declare their lower-bound limits are active, but agent i declares the highest bid, i.e., $|\underline{A}_i(\tau)| = N$ and $v(\pi_i(\tau)) \leq v(\pi_j(\tau))$ for all $j \neq i$, ties are broken assuming an ordering over agents;*
$\mathcal{C}_4$. *agent i finds a fixed point local solution and none of the other agents will declare part of their unit constraints are active, i.e., $\pi_i(\tau) = \pi_i(\tau + \delta\tau)$ and $\underline{A}_i(\tau + \delta\tau) = \overline{A}_i(\tau + \delta\tau) = \varnothing$, for all $\delta\tau \in \mathbb{Z}$.*

The total ordering among termination conditions is $\mathcal{C}_m > \mathcal{C}_{m+1}$, $\forall m \in \{1, 2, 3\}$.

Proof. We distinguish between the four termination conditions, for each of them we discuss corresponding scenarios.

$\mathcal{C}_1$. Assume $i \in \overline{A}_i(\tau)$. The following arguments hold for all four scenarios that match to this termination condition. If agent i declares its upper-bound limit is active, then it cannot

---

[4] Eight situations correspond to two possibility by set (whether or not a set is empty) raised at power of three sets, i.e., $2^3 = 8$.

generate power output more than its upper-bound limit. On the other hand, if it generates less than its maximum (declared) power output, the overall expected cost among all units will increase. This is mainly due to the way the DPGO algorithm proceeds. It divides load demand among units according to their local cost functions, larger power outputs are assigned to agents with weaker local cost functions. Hence, agent i best schedule is to produce to maximum power output (it declares). This proves the first termination condition is sufficient to state agent i is terminal.

$\mathcal{C}_2$. Assume $i \in \underline{A}_i(\tau)$, $\overline{A}_i(\tau) = \varnothing$ and $A_i(\tau) \neq \varnothing$. In this condition, it only remains one situation. This situation occurs when both $A_i(\tau)$ and $\underline{A}_i(\tau)$ are not empty sets. If unit i declares its lower-bound limit is active, then one can set its power output to be the minimum possible and make the other generator units to supply the remaining load. To better understand the rational behind this schedule, notice the following cases. On the one hand, if unit i generates more power output, then the overall expected cost of the power generation system will increase. On the other hand, if unit i generates the minimum, then the other units are guaranteed to supply the remaining load demand on their own and preserve a lower overall expected cost. This demonstrates the rational of the second termination condition.

$\mathcal{C}_3$. Assume $|\underline{A}_i(\tau)| = N$ and $v(\pi_i(\tau)) \leq v(\pi_j(\tau))$ for all $j \neq i$. There is only one situation that matches to this termination condition — including scenarios where all units declare their lower-bound limits are active. In that case, we set power output of the unit with the lowest bid at the maximum power output it can, and claim this unit is terminal. Clearly, this is a heuristic strategy, which ensures the expected cost will not increase as only the agent — say i — with the minimum bid terminates. Furthermore, we are guaranteed the other units can supply the remaining load demand. This is mainly because, before applying this heuristic strategy, units already met the load demand while agent i were producing power output less than what is required by the heuristic strategy, i.e., the maximum possible.

$\mathcal{C}_4$. Now assume $\pi_i(\tau) = \pi_i(\tau + \delta\tau)$ and $\underline{A}_i(\tau + \delta\tau) = \overline{A}_i(\tau + \delta\tau) = \varnothing$, $\forall \delta\tau \in \mathbb{Z}$. This condition corresponds to situations where all agents satisfy their unit constraints and have found a fixed local solution. This corresponds to the optimum solution of the optimization problem with only equality constraints.

This ends the proof.□

It is worth noticing that the number of agents involved in the planning process will decrease as agents become terminal. In fact, agents that are declared terminal at round τ are removed from succeeding sets $\underline{A}_i(\tau + \delta\tau)$, $\overline{A}_i(\tau + \delta\tau)$ and $A_i(\tau + \delta\tau)$, for any $\delta\tau \in \mathbb{Z}$. In addition, the remaining agents reset the optimization problem by removing from the load demand the amount of power the terminal agents will generate.

Termination conditions often involve local decision variables to be fixed according to unit constraints, which may violate the load and generation balance. Such a situation is referred to as a conflicting situation. Specifically, termination conditions $\mathcal{C}_1$ and $\mathcal{C}_2$ violate the prescription of the DPGO algorithm to satisfy unit constraints, which makes the overall schedule unable to satisfy the load demand. To solve these conflicts we rely on messages of different types, each of which corresponds to either termination conditions or normal situations.

5.2. Conflict protocols

In this subsection, we present communication protocols that agents use throughout the distributed planning process, in order to communicate with one another. In particular, we discuss

messages that agents exchange to handle conflicting situations. To handle conflicts, non-terminal agents reset the problem whenever a conflict occurs. The new problem consists in decision variables that belong to non-terminal agents, and the remaining load demand. Then, these agents restart the distributed planning process with that novel problem, and the process continues until all agents are terminal according to Theorem 1. The termination of the algorithm is guaranteed by observing that each conflicting situations occurs only after an agent is declared to be terminal, hence there is a finite number of conflicting situations to be handled. Overall, we distinguish between two types of messages including normal and terminal messages.

A message is defined as a four tuple $\langle type, policy, signal, bid \rangle$, where the type of message is whether it is a normal or a terminal message. A message is said to normal whenever agent i is in none of the termination conditions (see Theorem 1) at the τ-th round of the algorithm. In that case, agent i can send a message referred to as $normal$, i.e., $\langle normal, \pi_i(\tau), \sigma_i(\tau), v(\pi_i(\tau)) \rangle$. A message is said to be terminal whenever agent i in one of the termination conditions (see Theorem 1) at the τ-th round of the algorithm. In that case, agent i can send a message referred to as $terminal$, i.e., $\langle terminal, \pi_i(\tau), \sigma_i(\tau), v(\pi_i(\tau)) \rangle$.

After the reception of normal and terminal messages at the τ–th round of the algorithm, agents update their local parameters, including local solution $\pi_i(\tau)$ and sets $A_i(\tau)$, $\underline{A}_i(\tau)$ and $\overline{A}_i(\tau)$. The updated solution $\pi_i(\tau)$ consists of local solutions $(\pi_i(\tau))_{i \in \{1,2,\ldots,N\}}$ agents share with one another. Next, the algorithm checks the consistency of the updated local solution with respect to system constraints, including load and generation balance. If the latter is not satisfied, then a conflicting situation occurs. Similarly, each agent includes the other agents and himself to sets $A_i(\tau)$, $\underline{A}_i(\tau)$ and $\overline{A}_i(\tau)$ according to signals received.

5.3. The P-DPGO algorithm

The P-DPGO Algorithm 1 is an offline and multi-agent distributed planning algorithm, in which each agent produces its local solutions by means of the DPGO algorithm and communicates with the other agents through the aforementioned communication protocols to handle conflicting situations.
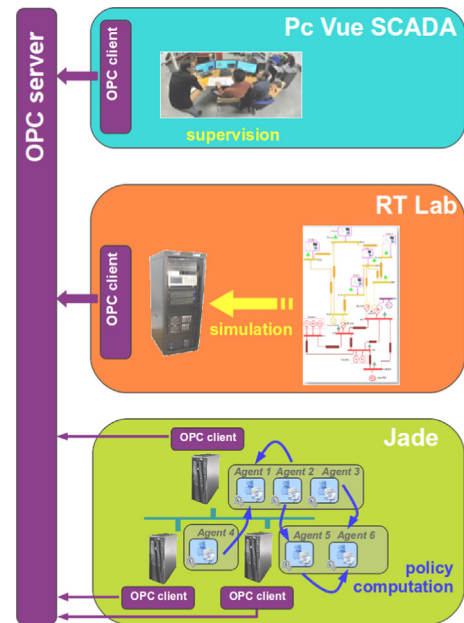

Fig. 1. Experimental platform.

Each agent starts with an initial local solution (see lines 2 and 3, Algorithm 1). A simple rule-of-thumb for the INITIALIZATION step is to spread the demand uniformly among available agents conditional on contingencies. Notice that initial local solutions may violate unit constraints.

Then, at each round of the P-DPGO algorithm, each agent checks whether its current local solution holds conflicts. If a conflict is detected, the P-DPGO algorithm re-initializes the current local solution by taking into account only non-terminal agents still involved in the planning process (see lines 5 and 6, Algorithm 1). As previously mentioned, the P-DPGO algorithm cannot cycle indefinitely since conflicting situations occur only after an agent has been declared terminal, that is in the worst case there is only N conflicts to handle.

Next, each agent computes its local solution based on the current one via the DPGO algorithm (see line 7, Algorithm 1). Each agent also keeps track of the signal and bid associated with the local solution just calculated. Recall that the DPGO algorithm cannot handle unit constraints, so the calculated local solution may violate these constraints. To tackle this limitation, the P-DPGO algorithm distinguishes between five cases. The first four cases correspond to terminal conditions, and the fifth case the remaining scenarios.

For each case, each agent adapts the local solution prescribed by the DPGO algorithm to handle its unit constraints. For instance, if the local solution prescribed by the DPGO algorithm violates the upper-bound limits, the P-DPGO algorithm replaces the power output prescribed by the maximum possible and sets the message accordingly (see subroutine SENDMESSAGE).

Finally, the agent wait for the messages the other agents have sent, and merge together these messages to update their local solution with that from the other agents in subroutine WAITMESSAGE. The P-DPGO algorithm repeats this procedure until the agent is declared to be terminal. The agent ends with the last local solution it communicated to the other agents before termination.

agents involved in the planning process) if the third termination condition never occurs. Indeed, we use an optimal strategy to assign values to decision variables for all termination conditions, except the third one. In practice, however, we note that its performances are equivalent and often higher than that of exact centralized solution methods. For this reasons, we are investigating conditions that would ensure the optimality of the solution returned by the P-DPGO algorithm within the context of smart-grid scenarios, and more other scenarios. Next, we demonstrate that the P-DPGO algorithm terminates in a finite number of iterations.

Theorem 2. *The P-DPGO algorithm is guaranteed to terminate after a finite number of iterations.*

Proof. As the algorithm proceeds, agents alternate between termination and normal conditions. Each agent can only experiment N termination messages, each of which correspond to a termination message of a single agent. During the remaining rounds of the algorithm, the P-DPGO algorithm restricts to the DPGO algorithm, which is guaranteed to terminate after a finite number of iterations. □

6. Experiments

To evaluate the performance the P-DPGO algorithm, we use an experimental test platform dedicated to "distributed energy". The main objective of this platform is to study electrical networks in the presence of decentralized generation. It is composed of a simulator and several generation systems, storages, loads, super capacitors, etc. Using this experimental platform, different approaches (especially multi-agent ones) for the coordination of generation systems of different types can be experimented.

Algorithm 1. The protocol based distributed projected gradient-descent based optimization (P-DPGO) algorithm.

```
 1  function P-DPGO(agent i)
 2  |   A_i(0) = ∅, Ā_i(0) = ∅ and A_i(0) = {1, 2, …, N}
 3  |   τ ← 0 and π_i(τ) ← INITIALIZATION(A_i(τ), Ā_i(τ), A_i(τ))
 4  |   while agent i is non-terminal do
 5  |   |   if ISCONFLICTING(π_i(τ)) then
 6  |   |   |   π_i(τ) ← INITIALIZATION(A_i(τ), Ā_i(τ), A_i(τ))
 7  |   |   ⟨π_i(τ+1), σ_i(τ+1), v(π_i(τ+1))⟩ ← DPGO(π_i(τ))
 8  |   |   SENDMESSAGE(π_i(τ+1), σ_i(τ+1), v(π_i(τ+1)))
 9  |   |   WAITMESSAGE(π_i(τ+1), A_i(τ+1), Ā_i(τ+1), A_i(τ+1))
10  |   |   τ ← τ + 1

11  function SENDMESSAGE(π_i(τ), σ_i(τ), v(π_i(τ)))
12  |   if 𝒞_1 then
13  |   |   return ⟨terminal, π_i^max(τ), σ_i(τ), v(π_i^max(τ))⟩
14  |   else if 𝒞_2 or 𝒞_3 then
15  |   |   return ⟨terminal, π_i^min(τ), σ_i(τ), v(π_i^min(τ))⟩
16  |   else if 𝒞_4 then
17  |   |   return ⟨terminal, π_i(τ), σ_i(τ), v(π_i(τ))⟩
18  |   return ⟨normal, π_i(τ), σ_i(τ), v(π_i(τ))⟩
```

5.4. Theoretical properties

This subsection discusses theoretical properties of the P-DPGO algorithm including performance guarantees and termination of the algorithm.

Using the P-DPGO algorithm, the agents terminate with an optimal solution (with respect to available information to all

6.1. Description of the experimental platform

The platform includes various tools and sub-platforms interacting together (Fig. 1). The core of the platform is RT-Lab (Real-Time Digital Simulation and Control LABoratory (Dufour et al., 2005): a distributed simulator that is able to simulate, in real-time, dynamic models of any parts of a power grid. RT-Lab allows us to
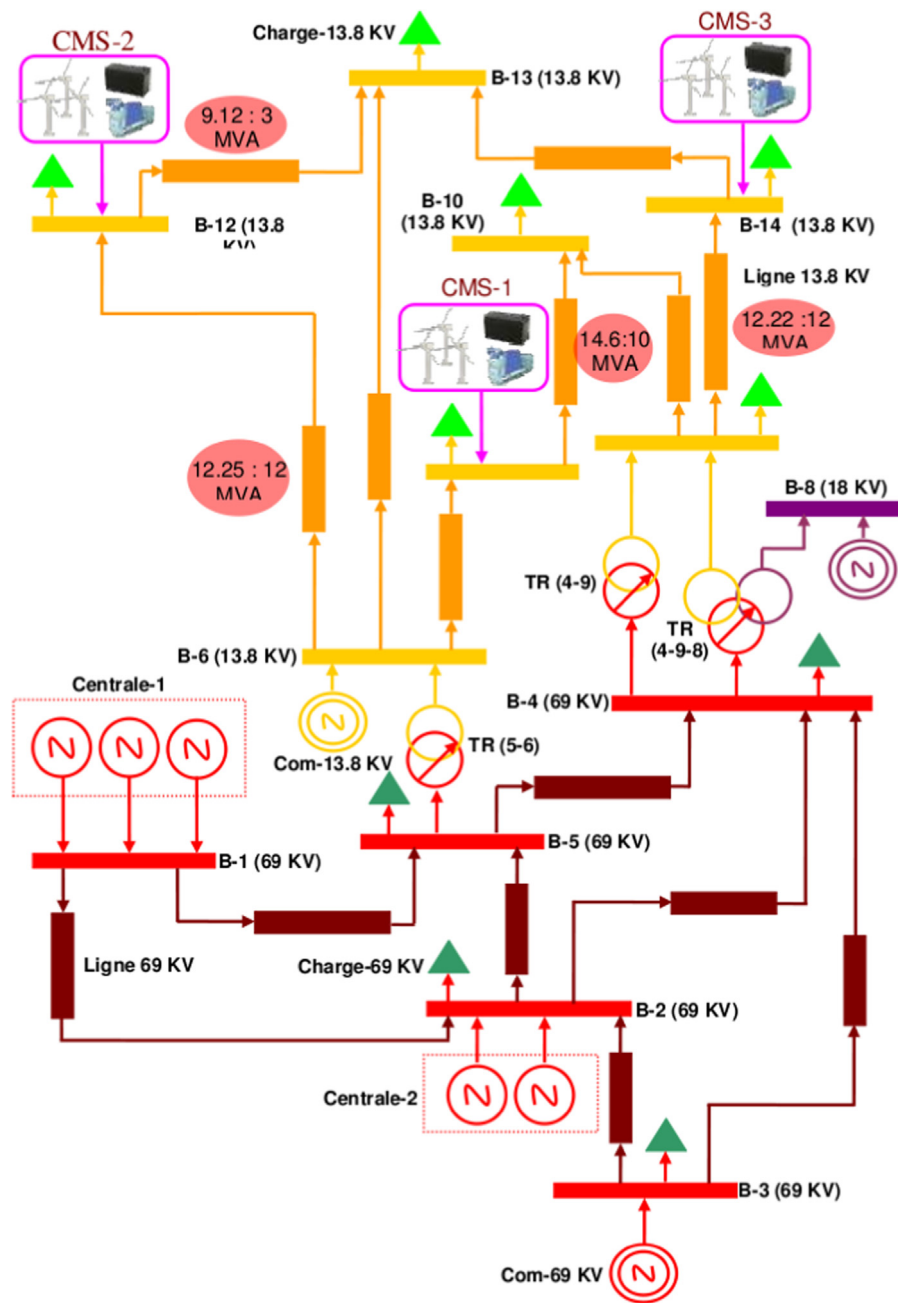
Fig. 2. Test network.

Table 2
Diesel power generators features.

p_{max} (kW)	Fuel consumption depending of the output profile p_i	t_{UP}	t_{DN}	t_{CL}
10.8	$5.13p_i^2 - 10.19p_i + 29.53$	80	80	80
6.30	$5.00p_i^2 - 10.00p_i + 29.72$	80	80	80
5.40	$4.94p_i^2 - 9.92p_i + 29.79$	80	80	80

test and design controllers for a large variety of devices including power converters, controllable generating units, renewable energy systems (wind turbine and photovoltaic panels). The interface with real equipments is possible through power amplifiers. The supervision of the network simulated with RT-LAb is based on PcVue which allows us to collect data and visualizes the dynamics of the simulation.

The agents we used to control generating units are implemented using Jade (Java Agent DEvelopment Framework) (Bellifemine et al., 2007). The multi-agent system is hosted by several PC connected by an ethernet network. Each PC can host several agents. Jade eases communications between agents through a support for message-passing: agents hosted in same PC communicate through Java RMI, otherwise HTTP protocol is used.

Jade, RT-Lab and PcVue are able to interact thanks to an OPC[5] server. Each of them used an OPC client to exchange data. For Jade, each agent used its own OPC client with respect to distribution and decentralization properties (Fig. 1).

[5] The acronym 'OPC' stands for 'OLE' (object linking and embedding) for process control.

Fig. 3. Demand of power.

6.2. Experimental network

For our experiments, we use an IEEE 14 nodes network slightly modified as depicted in Fig. 2. Several virtual power plants have been added whose controllable generating units are diesel power generators (named CMS in Fig. 2). Each of them is able to generate profiles p_i that range from 10% to 95% of $p_{i\max}$, the maximum capacity of production, i.e.,

$$0(C1) : 0.10*p_{i\max} \le p_i \le 0.95*p_{i\max}. \tag{2}$$

In addition, the fuel consumption of each generating unit is quadratic. It is worth noticing that the units with the biggest production capacity are also the units with the highest cost. Finally, parameters t_{UP}, t_{DN} and t_{CL} are operational time limits (in minutes) corresponding respectively to the minimum up time, the minimum down time and the cold start time. For each scenario, we consider the demand over 24 h.

It is worth noticing that the main source of uncertainty in this network is the unpredictable breakdowns of diesel power generators. To address this aspect, we add all possible breakdown situations as part of the observations the agents can receive. Specifically, we use a Poisson distribution to describe the probability that a breakdown occurs for each generator unit over periods of time, independently from the other generator units.

The fuel consumption functions (Table 2) are of crucial importance in competitive settings. For instance, under the assumption of rational agents, this information can be used to predict the schedules of the other agents, which is undesirable in a competitive setting. Hence, we would like this information to be preserved during information exchange between agents. This highlights the necessity for a distributed algorithm that can preserve private information of each agent.

Each agent in charge of a diesel power generator starts its life cycle by registering on the Jade platform. In this way, its name and IP address are stored by a Directory Facilitator agent, which is in charge of the 'Yellow pages' phone book. Then, agents can communicate with one another through messages. Once all agents are registered, they all receive at a DSUC instance corresponding to one specific scenario from a human operator interacting with the RT-Lab platform.

6.3. Experimental results

This subsection presents the experimental results obtained for different scenarios. We distinguish between four scenarios along with the associated DSUCs. The first set of experiments (*Scenario* 1) is our baseline since it assumes no unpredictable event (e.g., breakdowns) can occur at the execution phase. The second set of experiments (*Scenario* 2) extends Scenario 1 and allows unpredictable breakdowns to occur. These experiments highlight the necessity of representing unit commitment problems using DSUCs. The third set of experiments (*Scenario* 3) compares performances we obtain using a reactive approach, i.e., considering a planning horizon $T=1$ (this approach is referred to as P-DPGO-one-stage) versus an approach that takes into account long term effects of present decisions, i.e., the P-DPGO algorithm. The last set of experiments (*Scenario* 4) provides the opportunity to evaluate the scalability of the P-DPGO algorithm as the number of agents increases.

Demand profiles we use for each of these scenarios are depicted in Fig. 3. For each scenario, we run the P-DPGO algorithm, simulated the resulting policies and reported: (a) the number of units involved; (b) load profiles for all units involved; (c) breakdown events when they occur. In Scenarios 3 and 2, we also reported the cumulated cost over the planning horizon each experimented approach yields. In Scenario 4, we reported the total time required to find a solution using the P-DPGO algorithm as the number of agents increases.

Scenario 1. In this scenario, we consider two configurations with five and six diesel power generators, respectively. We simulate policies the P-DPGO algorithm pre-computed in face of different demand profiles over 24 h. For the two configurations, Figs. 4 and 5 reported traces of these simulations. Cumulative histograms allow us to easily visualize the commitment of each unit. Clearly, other all tested instances (as demonstrated by Figs. 4 and 5), P-DPGO always produces policies that can meet the demand in the absence of unpredictable events.

Scenario 2. Here, we extend Scenario 1 and simulate unpredictable breakdowns of the diesel power generators. It is in such a setting that P-DPGO fully demonstrates its advantage over classical (centralized) approaches. It produces (in a distributed manner) policies that can take into account events that may occur

**Fig. 4.** Scenario 1: DSUC with 5 units.

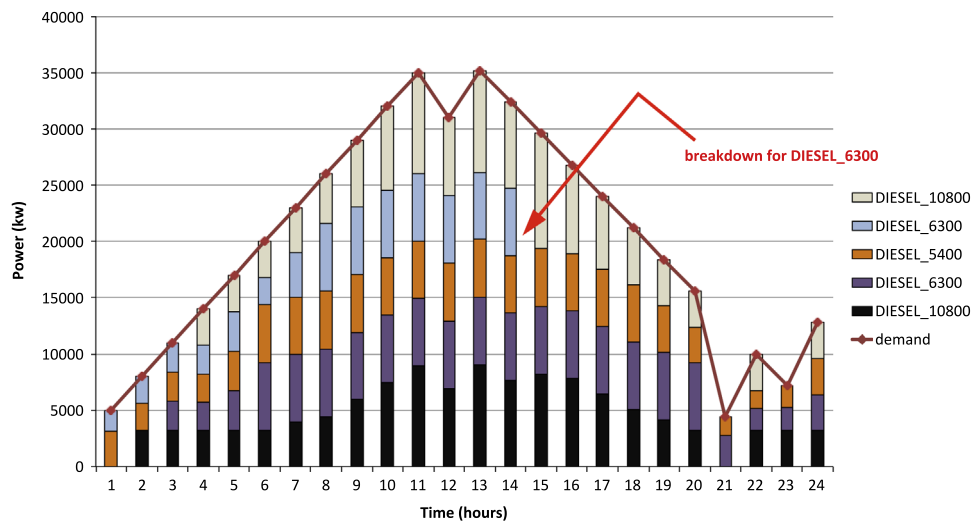**Fig. 5.** Scenario 1: DSUC with 6 units.

Fig. 6. Scenario 2 with 5 units and 1 breakdown occurring.

at the execution phase, e.g., unpredictable breakdowns as illustrated in Figs. 6 and 7. Specifically, Fig. 6 shows that P-DPGO produces policies that can automatically switch to situations where one unit (here unit 'DIESEL_6300') is no longer available for the remaining of the planning horizon while still meeting the demand. Fig. 7 presents another extreme case where many units
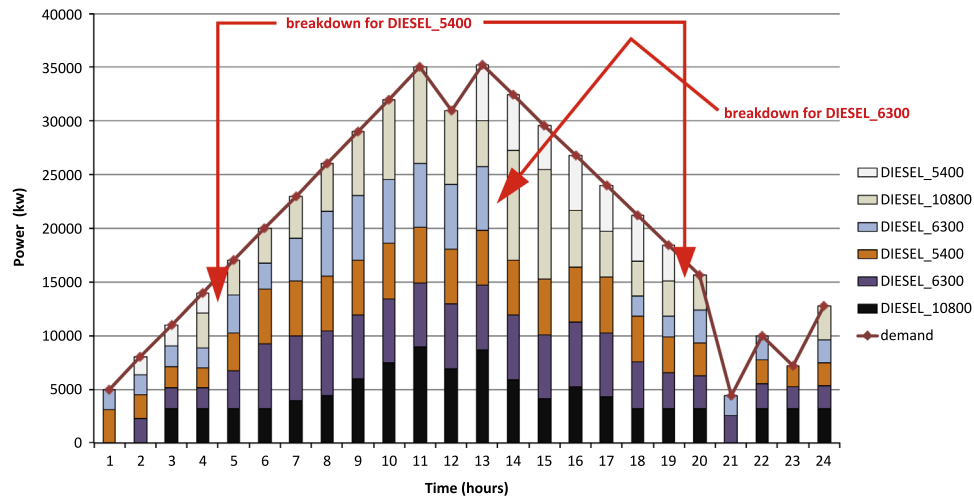
Fig. 7. Scenario 2 with 6 units and 3 breakdowns occurring.

Fig. 8. Cumulated cost for Scenario 2 with 3 breakdowns.

Fig. 9. Scenario 3: DSUC with 3 units solved with P-DPGO.

(here units 'DIESEL_5400' and 'DIESEL_6300') experienced break-downs at different period in the planning horizon. Yet, P-DPGO preserves ability to meet the demand based on the remaining units. Non surprisingly, Fig. 8 shows that these breakdowns can significantly affect the total costs of P-DPGO policies. Specifically,

when unpredictable events may occur, the total cost of P-DPGO policies is about 2.5% of than the case when no unforeseen events may happen.

*Scenario* 3. This scenario compares the performances of P-DPGO-one-stage and P-DPGO algorithms. P-DPGO-one-stage

**Fig. 10.** Scenario 3: DSUC with 3 units solved with P-DPGO-one-stage.

**Fig. 11.** Scenario 3: DSUC cumulated cost difference between P-DPGO-one-stage and P-DPGO.

Fig. 12. Scenario 3: DSUC with 3 units solved with P-DPGO.

refers to the use of P-DPGO with a planning horizon $T=1$, in other words without taking into account long term effects of present decisions. To compare P-DPGO-one-stage and P-DPGO

performances, we consider three diesel power generator units. Figs. 9 and 10 report demand and production profiles over 24 h for P-DPGO and P-DPGO-one-stage algorithms, respectively. For this

**Fig. 13.** Scenario 3: DSUC with 3 units solved with P-DPGO-one-stage.

Fig. 14. Time computing with different DSUC sizes.

specific configuration, both approaches meet the demand over the entire planning horizon. However, Fig. 11 offers the opportunity to compare the qualitative performances of each approach with respect to the total costs they induce. In particular, we note that P-DPGO-one-stage requires a total cost that is 7% higher than that from P-DPGO. Clearly, the advantage of the P-DPGO algorithm comes from its ability to anticipate future effects of present choices, whereas P-DPGO-one-stage cannot. For example, in Figs. 9 and 10, we observe that P-DPGO allows to turn off the 'DIESEL_10800' (the unit with the highest cost) more often than with P-DPGO-one-stage. In addition there are many cases in which P-DPGO-one-stage fails to meet the demand as it cannot take into account future effects of present decisions. To illustrate this point, we slightly modify the demand previously used: at time $t=16$ we increase the demand from 8000 to 15 000 and at time $t=17$ from 15 000 to 20 000 (demand profile 2 in Fig. 3).

In this new case, P-DPGO-one-stage fails to satisfy the load demand. Indeed, at time $t=15$, P-DPGO-one-stage turns off the 'DIESEL_10800' as illustrated in Fig. 13. Due to the operation time limit of the unit (Table 2), unit 'DIESEL_10800' cannot be restarted and fails to meet the load demand at $t=16$. In contrary, the P-DPGO algorithm uses a long term planning, hence it can handle this new case (Fig. 12).

Scenario 4. The complexity of DSUCs increases linearly with the planning horizon and exponentially with the number of generator units. For this reason, in demonstrating the scalability of the

P-DPGO algorithm, we run experiments for increasing the number of units. To this end, over each tested configuration, we keep track of both the number of messages agents exchanged during the planning phase, and the computational time required in order to find the final solution. Overall, both the computational time and the number of messages grow exponentially with increasing number of units. For instance, the number exchanged messages we recorded is about 85 000 for three units and up to 5 millions for six units. Similar results hold for the computation time as the number of units increases (Fig. 14).

7. Conclusion

This paper explores the economic dispatch and unit commitment (UC) problem which consists in finding the optimal schedules and amounts of power to be generated by a set of distributed units in response to an electricity demand. The motivation of this work is to propose a distributed approach to solve UC problems taking into account uncertainty from both demand and supply and privacy-preserving constraints.

We introduce an algorithmic framework that extends both stochastic and distributed approaches to UCs in order to ensure that electrical utilities do not explicitly communicate private information during the planning phase. We show how to recast distributed stochastic unit commitment problem into linearly

constrained quadratic programs. The main contribution of the paper concerns an offline and multi-agent distributed planning algorithm (termed P-DPGO) which combines a projected gradient descent and a communication protocol that allows electrical utilities to collectively find an optimal schedule strategy while ensuring that private information remain unshared.

The performance of P-DPGO is evaluated over different stochastic scenarios using a real-time experimental platform. The experiments highlight the salient features of P-DPGO: it preserves private information, it takes into account long term decision effects and it handles uncertainty in supply and demand.

In the future, we plan to explore applying this approach to competitive scenarios, with self-interested agents. Furthermore, we will investigate P-DPGO theoretical guarantees that can ensure that we end with an optimal policy.

References

Aoki, K., Itoh, M., Satoh, T., Nara, K., Kanezashi, M., 1989. Optimal long-term unit commitment in large scale systems including fuel constrained thermal and pumped-storage hydro. IEEE Trans. Power Syst. 4 (3), 1065–1073.

Bellifemine, F.L., Caire, G., Greenwood, D., 2007. Developing Multi-Agent Systems with JADE. John Wiley and Sons.

Borghetti, A., Frangioni, A., Lacalandra, F., Lodi, A., Martello, S., Nucci, C., Trebbi, A., 2001. Lagrangian relaxation and tabu search approaches for the unit commitment problem. In: 2001 IEEE Porto Power Tech Proceedings, vol. 3.

Bouffard, F., Galiana, F., Conejo, A., 2005. Market-clearing with stochastic security—part i: formulation. IEEE Trans. Power Syst. 20 (4), 1818–1826.

Boyd, S., Vandenberghe, L., 2004. Convex Optimization. Cambridge University Press, New York, NY, USA.

Cohen, A., Yoshimura, M., February 1983. A branch-and-bound algorithm for unit commitment. IEEE Trans. Power Appar. Syst. PAS 102 (2), 444–451.

Dechter, R., 2003. Constraint optimization, Constraint Processing. Morgan Kaufmann, San Francisco, pp. 363–397.

Dufour, C., Abourida, S., Belanger, J., 2005. Hardware-in-the-loop simulation of power drives with rt-lab. In: Proceedings of the 6th International IEEE Conference on Power Electronics and Drives Systems, PEDS 2005, vol. 2, November, pp. 1646–1651.

Fisher, M., 2004. The Lagrangian relaxation method for solving integer programming problems. Manage. Sci. 50, 1861–1871.

Guan, X., Zhai, Q., Papalexopoulos, A., 2003. Optimization based methods for unit commitment: Lagrangian relaxation versus general mixed integer programming. In: Power Engineering Society General Meeting, vol. 2, July 2003, IEEE, Toronto, Canada.

Kim, B., Baldick, R., 1997. Coarse-grained distributed optimal power flow. IEEE Trans. Power Syst. 12 (2), 932–939.

Kok, J., Scheepers, M., Kamphuis, I., 2010. Intelligence in electricity networks for embedding renewables and distributed generation. In: Intelligent Infrastructures. Intelligent Systems, Control and Automation: Science and Engineering. Springer. pp. 179–209.

Kumar, A., Faltings, B., Adrian, P., 2009. Distributed constraint optimization with structured resource constraints. In: Proceedings of the 8th International Conference on Autonomous Agents and Multiagent Systems, AAMAS 2009, vol. 2, pp. 923–930.

Lesser, V., Tambe, M., Ortiz, C.L. (Eds.), 2003. Distributed Sensor Networks: A Multiagent Perspective. Kluwer Academic Publishers, Norwell, MA, USA.

Mathews, G., Durrant-Whyte, H., Prokopenko, M., 2009. Decentralised decision making in heterogeneous teams using anonymous optimisation. Robot. Auton. Syst. 57 (3), 310–320.

Miller, S., Ramchurn, S., Rogers, A., 2012. Optimal decentralised dispatch of embedded generation in the smart grid. In: Proceedings of the 11th International Conference on Autonomous Agents and Multiagent Systems, AAMAS 2012, vol. 1, pp. 281–288.

Modi, P.J., Shen, W.-M., Tambe, M., Yokoo, M., 2005. Adopt: asynchronous distributed constraint optimization with quality guarantees. Artif. Intell. 161 (1–2), 149–180.

Nikovski, D., Zhang, W., 2010. Factored Markov decision process models for stochastic unit commitment. In: 2010 IEEE Conference on Innovative Technologies for an Efficient and Reliable Electricity Supply (CITRES), pp. 28–35.

Nowak, M., Rmisch, W., 2000. Stochastic lagrangian relaxation applied to power scheduling in a hydro-thermal system under uncertainty. Ann. Oper. Res. 100 (1–4), 251–272.

Ramchurn, S., Vytelingum, P., Rogers, A., Jennings, N., 2012. Putting the 'smarts' into the smart grid: a grand challenge for artificial intelligence. Commun. ACM 55 (4), 86–97.

Sheble, G., Fahd, G., 1994. Unit commitment literature synopsis. IEEE Trans. Power Syst. 9 (1), 128–135.

Snyder, W., Powell, H., Rayburn, J., 1987. Dynamic programming approach to unit commitment. IEEE Trans. Power Syst. 2 (2), 339–348.

Takriti, S., Birge, J., Long, E., 1996. A stochastic model for the unit commitment problem. IEEE Trans. Power Syst. 11 (3), 1497–1508.

Tarim, A., Manandhar, S., Walsh, T., 2006. Stochastic constraint programming: a scenario-based approach. Constraints 11 (1), 53–80.

Tsitsiklis, J., Bertsekas, D., Athans, M., 1986. Distributed asynchronous deterministic and stochastic gradient optimization algorithms. IEEE Trans. Autom. Control 31 (9), 803–812.

Vytelingum, P., Ramchurn, S., Voice, T., Rogers, A., Jennings, N., 2010. Trading agents for the smart electricity grid. In: Proceedings of the 9th International Conference on Autonomous Agents and Multiagent Systems, AAMAS 2010, vol. 1, pp. 897–904.

Zhang, W., Wang, G., Xing, Z., Wittenburg, L., 2005. Distributed stochastic search and distributed breakout: properties, comparison and applications to constraint optimization problems in sensor networks. Artif. Intell. 161 (1–2), 55–87.

Zilberstein, S., Washington, R., Bernstein, D.S., Mouaddib, A.-I., 2002. Decision-theoretic control of planetary rovers, Revised Papers from the International Seminar on Advances in Plan-Based Control of Robotic Agents. Springer-Verlag, London, UK, pp. 270–289.